

ASEN 5519 - ALGORITHMIC MOTION PLANNING
HOMEWORK #3
JOE MICELI
23 OCT 2020

PROBLEM 1

- (a) Complete Planning Algorithm - An algorithm that is guaranteed to find a path if one exists. If a path does not exist, the complete algorithm will end in failure (it is guaranteed to not produce a false positive).
- (b) Optimal algorithm - An algorithm that produces the "best" solution (in this case, path) based on the metrics used to evaluate each solution. Examples of an optimum could be shortest ~~path~~ path, minimum energy required, ~~fastest~~ fastest time to find a path, etc.
- (c) A wavefront planner is complete only w.r.t. to the discretization used to obtain the wavefront. This is because discretization of a space will unavoidably cause some loss of the space. So if any portion of ~~all~~ the valid path(s) in a space are lost during discretization, the algorithm will not be complete.

Similarly, wavefront planners are not always optimal but could potentially be used to find an optimal path if the C-space it was operating in was designed to account for the metrics being used to evaluate an "optimal solution." An example

A wavefront planner is not an optimal planning algorithm. It is not guaranteed to produce the best solution if multiple exist. In these cases, ~~all~~ cells ~~can~~ have multiple equally valued neighbors and a wavefront planner randomly selects one of them. The resulting path lengths will be equal but this could have drastic real-world differences on the system (i.e. energy required by each node).

PROBLEM 2

(a)

$$\nabla U(q) = \nabla U_{\text{attr}}(q) + \nabla U_{\text{rep}}(q)$$

$$\nabla U_{\text{attr}} = \begin{cases} \zeta(q - q_{\text{goal}}) & , d(q, q_{\text{goal}}) \leq d^*_{\text{goal}} \\ \frac{d^*_{\text{goal}} \zeta(q - q_{\text{goal}})}{d(q, q_{\text{goal}})} & , d(q, q_{\text{goal}}) > d^*_{\text{goal}} \end{cases}$$

$$\nabla U_{\text{rep}} = \sum_i \nabla U_{\text{rep}_i}$$

$$\nabla U_{\text{rep}_i} = \begin{cases} \eta \left(\frac{1}{\alpha_i} - \frac{1}{d_i(q)} \right) \frac{1}{d_i^2(q)} \nabla d_i(q) & d_i(q) \leq Q_i^* \\ 0 & d_i(q) > Q_i^* \end{cases}$$

Algorithm:

while $\nabla U(q(i)) > \epsilon$:

$$q(i+1) = q(i) + \alpha(i) \nabla U(q(i))$$

$i = i + 1$

end while

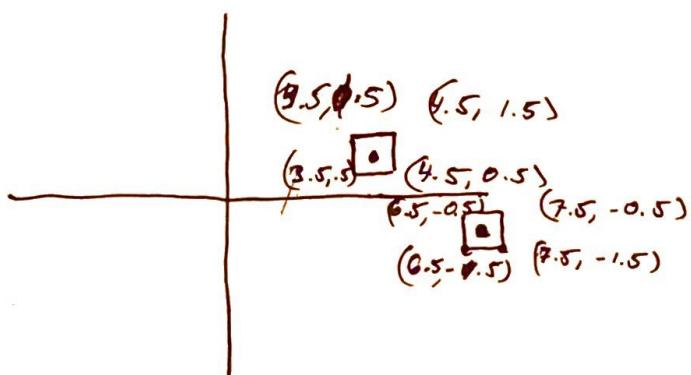
$$d_i(q) = \min_{c \in Q_0} d(q, c)$$

$$\nabla d_i(q) = \frac{q - c}{d(q, c)}$$

[Pseudo-Code]

Obstacles:

- Need to specify vertices
- Method of returning the point on the obstacle closest to current config, q :
 - ↳ Could use primitives to generate list of boundary points
 - ↳ Could check every boundary point for the closest one to current configuration
- Centered @ $(4, 1)$, lengths of 1



GetDistanceToObstacle(q)
 $\text{min_distance} = \text{GetDistance}(q, \text{boundary_points}[0])$
 for point in boundary-points

$\text{distance} = \text{GetDistance}(q, \text{point})$

if ($\text{distance} < \text{min_distance}$)

$\text{min_distance} = \text{distance}$

$\text{min_point} = C$

11) Get Boundary Points (

for ~~(~~ $\min_x \leq x \leq \max_x$)

boundary_points.push_back($x, P_1(x)$)

for ($\min_y \leq y \leq \max_y$)

b-p.push_back($P_2(y), y$)

for ($\max_x \leq x \leq \min_x$)

b-p.push_back($x, P_3(x)$)

for ($\min_y \leq y \leq \max_y$)

b-p.push_back($P_4(y), y$)

i) PLEASE SEE ATTACHED IMAGES/CODE

ii) I chose d^* goal as 5.0 which is a relatively large value for the workspace size but I wanted the attractive portion of the gradient to become large early on so obstacles ~~nearby~~ nearby would not cause the robot to get stuck.

Similarly, I wanted the robot to be "allowed" to get close to obstacles so I set Q_i^* to 1.0 for both, $i=1$ and $i=2$.

iii) PLEASE SEE ATTACHED IMAGES/CODE

iv) 9.87916 m is the length of the path.

v) No, different values for d^* goal and Q_i^* would almost always produce different paths & path lengths. These values impact how the robot behaves as it approaches an obstacle and how it avoids the obstacle. It can also produce local minima that prevent the robot from ever getting to the goal.

[IMPACT OF PARAMETERS]

$\epsilon = 0.25$ → Fixed in problem statement

$$\zeta = 0.1$$

$$\eta = 0.1$$

$$d_{goal} = 5.0$$

$$Q^+ = 1.0$$

} Gradient is 0 but robot stops
2.5m away from goal
↳ Local minimum!

$$\zeta = 2.0$$

$$\eta = 0.5$$

$$d_{goal} = 5.0$$

$$Q^+ = 1.0$$

} Robot is able to get close to
obstacles without getting stuck
because attractive gradient is strong
enough

(b)

- i) The values I settled on were obtained after dozens of trials. The algorithm is extremely sensitive to changes in any parameter: α , E , ϵ_0 , η , d_{goal} , Q^*

Ultimately, I was not able to find parameters that produced "collision free paths" for either workspace. My ~~initial~~ approach was to increase Q^* so that more obstacles would influence the gradient (even if it was far away). I chose a lower value for d_{goal} because I wanted obstacle avoidance to take precedence over achieving the goal. In workspace 1, this worked for avoiding obstacles but the robot became stuck in a local minimum.

In Workspace 2, the robot achieved the goal but collided w/ obstacles along the way. This planning problem was particularly challenging because the initial conditions produced a gradient descent direction that sent the robot straight into an obstacle.

[The exact values chosen for either workspace are listed in the README.txt provided with the code submission]

- ii.) PLEASE SEE ATTACHED IMAGES
- iii.) PLEASE SEE ATTACHED README
- iv.) No, both of these planning problems were very sensitive to changes in d_{goal} and Q^* . Slight modifications to either parameter drastically changes the potential function and therefore, the path produced from a gradient descent algorithm.

PROBLEM 3

- (a) PLEASE SEE ATTACHED CODE / IMAGES
- (b) Path Length of Workspace 1: 19.75
Path Length of Workspace 2: 43.75
- (c) As grid size gets smaller, I would expect the path length to decrease but converge. If cells of equal value are selected ~~at random~~, randomly with equal probability, I would expect path length to converge towards the ~~minimum~~ minimum path length that can physically be achieved in the workspace.
- (d) This algorithm performed much better than the gradient descent algorithm in these workspaces. In fact, the gradient descent algorithm was not able to find a collision free path ~~at all~~ (I was not able to find gradient descent parameters that produced a collision free path).
- Both algorithms require discretization of the full workspace and therefore require similar computational effort. The wavefront planner was more difficult to implement for me ~~but given the results, I would choose this algorithm over~~ but given the results, I would choose this algorithm over gradient descent.

[PSUEDO CODE]

Class : Cell:

Public:

Position

Value → Initialize as 0

std::vector<cell> neighbors(4)

[NOTE: Ended up not using this class]

- ① Discretize Workspace () → returns array or vector of all cells in workspace
 - ↳ Create cells
 - ↳ Initialize cell values w/ zero
 - ↳ Set cell neighbors
($x, y+step$)

NOTE: index will not correspond to cell value!

~~cell.neighbors.push_back(cell(x-step, y))~~

• • •
 ($x-step, y$) (x, y) ($x+step, y$)

•
 ($x, y-step$)

↳ Set goal cell

IF ($\text{Get Distance}(\text{cell}, \text{goal}) < 0.00001$):

$\text{cell.value} = 2$ → This should be the 1st cell in the vector of cells

[NOTE: This method did not work, I believe due to an issue passing pointers into function. Next time, try returning an object or learn how to use pointers]

Discretize Workspace \rightarrow Using a matrix/queue

if
x_{max} # x vals = $\frac{x_{\text{max}} - x_{\text{min}}}{\text{grid spacing}}$
x_{min}
y_{max} # y vals = $\frac{y_{\text{max}} - y_{\text{min}}}{\text{grid spacing}}$
y_{min}

for x in #xvals:

for y in #yvals:

$$\text{grid}(x)(y) = 0; \quad \begin{array}{l} \text{→ initialize every} \\ \text{location w/ 0} \end{array}$$

\hookrightarrow Location i in
C-space :

$$x_c = x_idx * \text{grid} + x_{\text{min}}$$

$$y_c = y_idx * \text{grid} + y_{\text{min}}$$

To get value @ x,y in C-space

$$\frac{x_c - x_{\text{min}}}{\text{grid spacing}} = x_idx$$

$$\frac{y_c - y_{\text{min}}}{\text{grid spacing}} = y_idx$$

$$\text{grid}(x_idx)(y_idx) = \text{value};$$

Label Collision Cells

```
for x in range grid-cells[x][]
    for y in range grid-cells[ ][y]
        check  $x_c = \dots$ 
         $y_c = \dots$ 
        if in collision( $x_c, y_c$ )
            grid-cells[x - idx][y - idy] = 1;
```

Label Remaining Cells

queue initialized by neighbors of g_{goal} & value of g_{goal}
 $g = \{(x+1, y, 2), (x, y+1, 2), (x-1, y, 2), (x, y-1, 2)\}$

while (queue not empty):

- get indices & values from queue ($q.front()$)
- remove this item from the queue
- check if indices have been explored
- if not, set their value in grid-cells to value from queue + 1
- Add its neighbors to the queue (if they ~~are in the workspace~~)

Find Path
while (cell-value > 2)

Start @ x_start, y_start

cell-value = grid-cells[x-start][y-start]

~~if~~ if (~~grid-cells[x-start+1][y-start]~~ == cell-value - 1)

cell-value = grid-cells[x-start+1][y-start] \rightarrow update value

~~x_idx =~~ $x_start + 1$, $y_idx = y_start$ \rightarrow update index
else if (...)

else if (...)

else if (...)

else
No valid neighbors